

# assignment-4

September 21, 2023

## 1 Grapes to Greatness: Machine Learning in Wine Quality

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
df=pd.read_csv('/content/winequality-red.csv')
df
```

```
[1]:      fixed acidity  volatile acidity  citric acid  residual sugar  chlorides \
0              7.4             0.700         0.00           1.9         0.076
1              7.8             0.880         0.00           2.6         0.098
2              7.8             0.760         0.04           2.3         0.092
3             11.2             0.280         0.56           1.9         0.075
4              7.4             0.700         0.00           1.9         0.076
...          ...          ...          ...          ...          ...
1594            6.2             0.600         0.08           2.0         0.090
1595            5.9             0.550         0.10           2.2         0.062
1596            6.3             0.510         0.13           2.3         0.076
1597            5.9             0.645         0.12           2.0         0.075
1598            6.0             0.310         0.47           3.6         0.067
```

```
      free sulfur dioxide  total sulfur dioxide  density  pH  sulphates \
0              11.0             34.0  0.99780  3.51         0.56
1              25.0             67.0  0.99680  3.20         0.68
2              15.0             54.0  0.99700  3.26         0.65
3              17.0             60.0  0.99800  3.16         0.58
4              11.0             34.0  0.99780  3.51         0.56
...          ...          ...          ...          ...          ...
1594            32.0             44.0  0.99490  3.45         0.58
1595            39.0             51.0  0.99512  3.52         0.76
1596            29.0             40.0  0.99574  3.42         0.75
1597            32.0             44.0  0.99547  3.57         0.71
1598            18.0             42.0  0.99549  3.39         0.66
```

```
      alcohol  quality
0          9.4         5
```

```

1          9.8          5
2          9.8          5
3          9.8          6
4          9.4          5
...
1594       10.5          5
1595       11.2          6
1596       11.0          6
1597       10.2          5
1598       11.0          6

```

[1599 rows x 12 columns]

```
[11]: df.describe()
```

```

[11]:      fixed acidity  volatile acidity  citric acid  residual sugar \
count    1599.000000      1599.000000  1599.000000      1599.000000
mean         8.319637         0.527821    0.270976         2.538806
std          1.741096         0.179060    0.194801         1.409928
min           4.600000         0.120000    0.000000         0.900000
25%           7.100000         0.390000    0.090000         1.900000
50%           7.900000         0.520000    0.260000         2.200000
75%           9.200000         0.640000    0.420000         2.600000
max          15.900000         1.580000    1.000000        15.500000

      chlorides  free sulfur dioxide  total sulfur dioxide      density \
count    1599.000000      1599.000000      1599.000000  1599.000000
mean         0.087467        15.874922        46.467792    0.996747
std          0.047065        10.460157        32.895324    0.001887
min           0.012000         1.000000         6.000000    0.990070
25%           0.070000         7.000000        22.000000    0.995600
50%           0.079000        14.000000        38.000000    0.996750
75%           0.090000        21.000000        62.000000    0.997835
max           0.611000        72.000000       289.000000    1.003690

      pH  sulphates  alcohol  quality
count    1599.000000  1599.000000  1599.000000  1599.000000
mean         3.311113    0.658149   10.422983    5.636023
std          0.154386    0.169507    1.065668    0.807569
min           2.740000    0.330000    8.400000    3.000000
25%           3.210000    0.550000    9.500000    5.000000
50%           3.310000    0.620000   10.200000    6.000000
75%           3.400000    0.730000   11.100000    6.000000
max           4.010000    2.000000   14.900000    8.000000

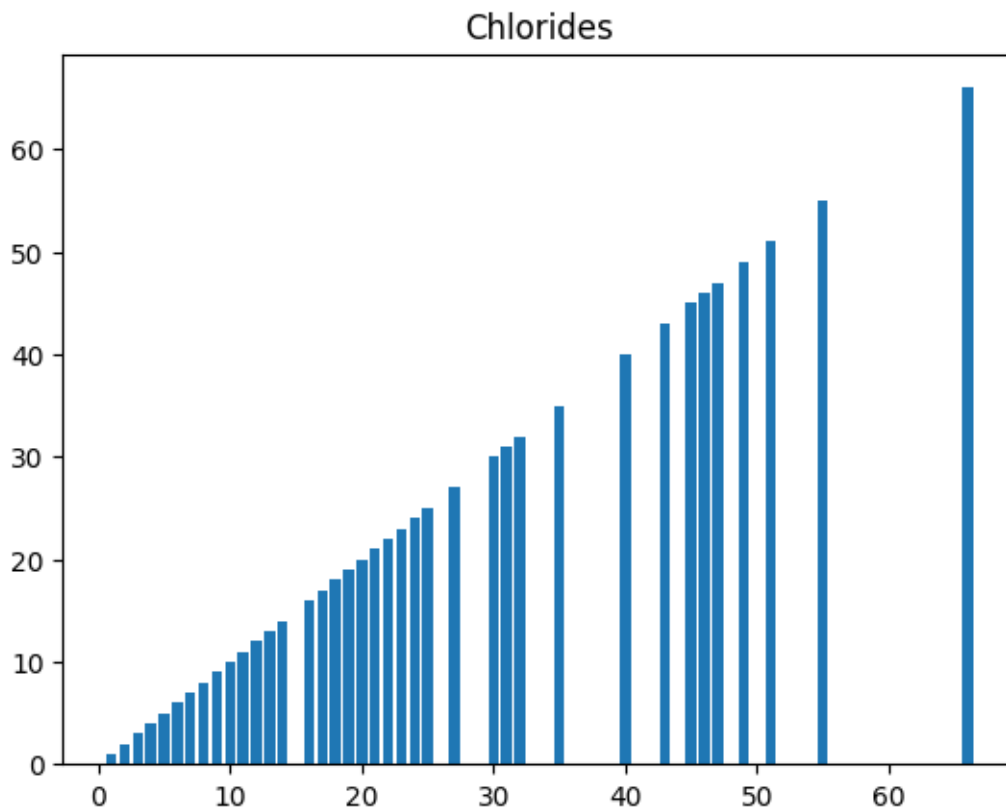
```

```


[35]: plt.bar(df.chlorides.value_counts(),df.chlorides.value_counts())
      plt.title("Chlorides")

```

```
plt.show()
```

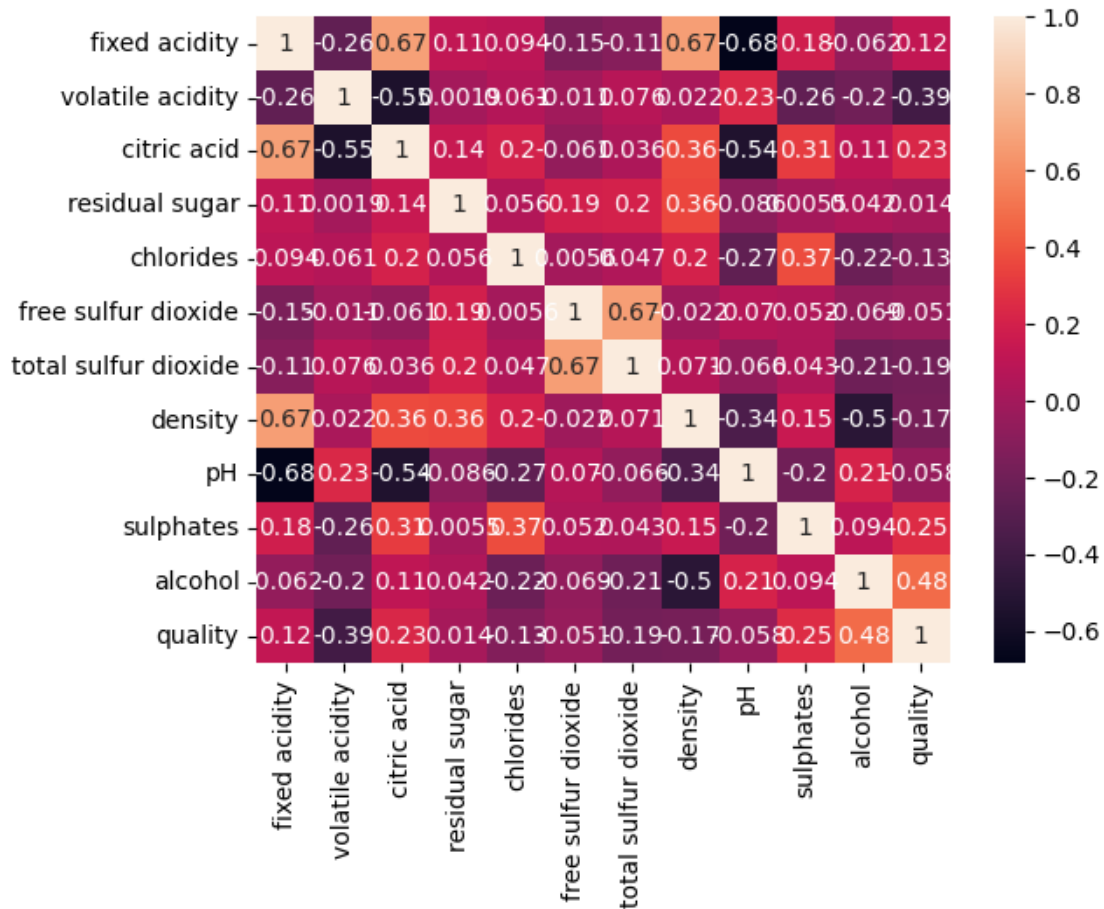


```
[7]: sns.jointplot(x="density",y="quality",data=df)  
plt.show()
```



```
[10]: sns.heatmap(df.corr(),annot=True)
```

```
[10]: <Axes: >
```



```
[14]: from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
X=df.drop("quality",axis=1)
y=df["quality"]
X.head()
```

```
[14]:    fixed acidity  volatile acidity  citric acid  residual sugar  chlorides  \
0           7.4           0.70           0.00           1.9       0.076
1           7.8           0.88           0.00           2.6       0.098
2           7.8           0.76           0.04           2.3       0.092
3          11.2           0.28           0.56           1.9       0.075
4           7.4           0.70           0.00           1.9       0.076

    free sulfur dioxide  total sulfur dioxide  density  pH  sulphates  \
0           11.0           34.0      0.9978  3.51       0.56
1           25.0           67.0      0.9968  3.20       0.68
2           15.0           54.0      0.9970  3.26       0.65
3           17.0           60.0      0.9980  3.16       0.58
```

4	11.0	34.0	0.9978	3.51	0.56
---	------	------	--------	------	------

	alcohol
0	9.4
1	9.8
2	9.8
3	9.8
4	9.4

```
[15]: sc=StandardScaler()
X_scaled=sc.fit_transform(X)
X_train,X_test,y_train,y_test=train_test_split(X_scaled,y,test_size=0.
↪2,random_state=42)
```

```
[15]: array([[ 0.21852997,  0.90601191,  0.20039205, ...,  1.09426457,
          0.48302886,  1.10483337],
        [-1.27524919, -1.77549685,  0.66254621, ..., -0.39596939,
          -0.40216729,  1.38643512],
        [ 1.48249695, -0.76993107,  1.02199944, ..., -0.07200549,
          0.54204194, -0.58477711],
        ...,
        [-0.6432657 ,  0.51495855, -1.08336951, ...,  1.28864292,
          -0.69723268, -0.86637886],
        [-0.24109439, -1.83136161,  0.4057939 , ...,  0.05758008,
          0.83710732,  1.38643512],
        [-1.44760832, -1.32857872, -0.05636026, ...,  0.51112954,
          -0.69723268,  2.8883111 ]])
```

```
[19]: from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix
le=LogisticRegression()
model=le.fit(X_train,y_train)
y_pred = model.predict(X_test)
cm = confusion_matrix(y_test, y_pred)
cm
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html>  
Please also refer to the documentation for alternative solver options:  
[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)  
n\_iter\_i = \_check\_optimize\_result(

```
[19]: array([[ 0,  0,  1,  0,  0,  0],
            [ 0,  1,  7,  2,  0,  0],
            [ 0,  0, 98, 32,  0,  0],
            [ 0,  0, 46, 76, 10,  0],
            [ 0,  0,  3, 30,  9,  0],
            [ 0,  0,  0,  1,  4,  0]])
```

```
[21]: from sklearn.metrics import accuracy_score, \
      ↪ confusion_matrix, classification_report, roc_auc_score, roc_curve
      accuracy_score(y_test, y_pred)
```

```
[21]: 0.575
```

```
[22]: pd.crosstab(y_test, y_pred)
```

```
[22]: col_0    4    5    6    7
      quality
3         0    1    0    0
4         1    7    2    0
5         0   98   32    0
6         0   46   76   10
7         0    3   30    9
8         0    0    1    4
```

```
[25]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
3	0.00	0.00	0.00	1
4	1.00	0.10	0.18	10
5	0.63	0.75	0.69	130
6	0.54	0.58	0.56	132
7	0.39	0.21	0.28	42
8	0.00	0.00	0.00	5
accuracy			0.57	320
macro avg	0.43	0.27	0.28	320
weighted avg	0.56	0.57	0.55	320

```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344:
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to
0.0 in labels with no predicted samples. Use `zero_division` parameter to
control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344:
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to
0.0 in labels with no predicted samples. Use `zero_division` parameter to
```

control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))  
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344:  
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to  
0.0 in labels with no predicted samples. Use `zero_division` parameter to  
control this behavior.  
_warn_prf(average, modifier, msg_start, len(result))
```

```
[27]: probability = model.predict_proba(X_test)[: ,1]  
probability
```

```
[27]: array([0.02483477, 0.02249377, 0.01574946, 0.01134614, 0.03202448,  
0.01491135, 0.00750669, 0.11853224, 0.02170618, 0.04960947,  
0.01348905, 0.13794541, 0.02885185, 0.02524655, 0.02042773,  
0.00977959, 0.03329229, 0.01903721, 0.00330975, 0.03185402,  
0.1797074 , 0.03375518, 0.06405695, 0.01701967, 0.07199606,  
0.03828524, 0.00477106, 0.08112635, 0.04058424, 0.01038622,  
0.03132 , 0.04077697, 0.01130467, 0.04878004, 0.02851514,  
0.02944521, 0.01913737, 0.06835138, 0.03249786, 0.02057124,  
0.02567618, 0.02042431, 0.006988 , 0.01425204, 0.01374079,  
0.03947643, 0.00374969, 0.01676378, 0.08833596, 0.07250037,  
0.01094354, 0.02734811, 0.09025171, 0.01177792, 0.04743786,  
0.01557644, 0.02994908, 0.09684315, 0.02995632, 0.06248571,  
0.01361311, 0.02148332, 0.03028016, 0.04898658, 0.00497826,  
0.01510166, 0.01183327, 0.04781076, 0.00504487, 0.01450206,  
0.00975868, 0.08957871, 0.0110487 , 0.03338539, 0.02254037,  
0.04057309, 0.00302335, 0.00481278, 0.02595905, 0.00514559,  
0.06328622, 0.00571601, 0.02323339, 0.03167214, 0.02807292,  
0.00498152, 0.01519544, 0.071952 , 0.00589653, 0.10965556,  
0.00579155, 0.04130017, 0.09166548, 0.0216257 , 0.00971574,  
0.01430115, 0.02018015, 0.02031068, 0.15743163, 0.02627518,  
0.1286683 , 0.03879046, 0.01325419, 0.10345292, 0.04510503,  
0.01106421, 0.01385973, 0.00778097, 0.03758517, 0.03038102,  
0.00631245, 0.00592356, 0.00255582, 0.0085207 , 0.02894913,  
0.01094736, 0.06563413, 0.01373076, 0.05589608, 0.01637058,  
0.01182063, 0.0395007 , 0.01413268, 0.02768522, 0.0678964 ,  
0.04112568, 0.00826422, 0.02335786, 0.03401157, 0.02896179,  
0.0110487 , 0.11028868, 0.03168882, 0.00694686, 0.07250037,  
0.06947725, 0.03621314, 0.00512094, 0.01887089, 0.0232298 ,  
0.01022772, 0.00711342, 0.00791129, 0.05265072, 0.01430848,  
0.03762969, 0.05080835, 0.13812178, 0.00655973, 0.01417962,  
0.01731361, 0.02659349, 0.03403913, 0.01742294, 0.0110487 ,  
0.00525754, 0.05574201, 0.01837955, 0.02863305, 0.02180492,  
0.00855826, 0.01296316, 0.00719941, 0.02473188, 0.02877727,  
0.01831045, 0.11902339, 0.03350339, 0.00606012, 0.01003063,  
0.01231073, 0.07909799, 0.01248025, 0.00938227, 0.00922779,  
0.00984898, 0.00563226, 0.01627412, 0.01257727, 0.26411334,
```



```

0.09242519, 0.00722218, 0.00895868, 0.0108839 , 0.09369835,
0.00256229, 0.14726728, 0.013673 , 0.00316726, 0.00938227,
0.05751076, 0.00648566, 0.01672627, 0.00407532, 0.01105994,
0.04465489, 0.13261131, 0.01265878, 0.02343221, 0.01106421,
0.04312645, 0.00825491, 0.02558483, 0.04303992, 0.00167707,
0.06365584, 0.02879044, 0.02108765, 0.02017862, 0.00209676,
0.04510175, 0.00560467, 0.01235811, 0.00154654, 0.01053313,
0.03851546, 0.07504775, 0.01853088, 0.04535837, 0.03899463,
0.0425178 , 0.00846753, 0.0104536 , 0.00717051, 0.23341551,
0.04884486, 0.04044068, 0.00566036, 0.00527932, 0.10618048,
0.00553059, 0.03591392, 0.01170368, 0.03609803, 0.00592857,
0.01106421, 0.06540582, 0.02807292, 0.06035659, 0.0120259 ,
0.04567327, 0.04198665, 0.06234175, 0.0061261 , 0.03219853,
0.02733513, 0.02068696, 0.04479066, 0.00227449, 0.00794547,
0.00523833, 0.0414681 , 0.04079487, 0.02742506, 0.35729187,
0.02987603, 0.02290243, 0.03998963, 0.01302718, 0.00323973,
0.01024649, 0.08956821, 0.01627412, 0.00917579, 0.04400462,
0.00749356, 0.03743595, 0.01220615, 0.00907912, 0.00546852,
0.0184154 , 0.00756899, 0.01566495, 0.01252444, 0.02138715,
0.04502192, 0.01317294, 0.00726284, 0.00533177, 0.05413836,
0.00522663, 0.12843439, 0.0386677 , 0.02689816, 0.01919488,
0.02910286, 0.03470307, 0.00711342, 0.03192086, 0.21076717,
0.0268413 , 0.01016235, 0.00992186, 0.01227231, 0.01608297,
0.00649988, 0.04338943, 0.00832624, 0.0088752 , 0.00587206,
0.00899913, 0.02178773, 0.00842123, 0.01201218, 0.04519739,
0.05589983, 0.01363781, 0.0050594 , 0.03334108, 0.00328764,
0.03285411, 0.03042666, 0.02367536, 0.04470886, 0.04878004,
0.0542435 , 0.02056979, 0.04453912, 0.00576608, 0.05107761])

```

```

[34]: model.predict([[7.4, 0.700, 0.00, 1.9, 0.076, 11.0, 34.0, 0.99780, 3.51, 0.56,
↪9.4]])

```

```

[34]: array([5])

```